

CADRE

Continual Adaptation for Driving with Robust Evolution

Problem Statement B5

Design a continual adaptation method for VLA autonomous driving models that incorporates new regional regulations, road layouts, and weather patterns while retaining prior driving competence across previously learned environments.

Datasets: BDD100K + nuScenes

Backbone: LLaVA-v1.5-7B (Frozen)

Method: EWC + Experience Replay + LoRA Adapters

Date: August 4, 2026

Key Results

BWT = -1.17% | FWT = +23.33% | Plasticity = 98.46%

Parameter Efficiency = 99.65% | CDAR = 0.9735

Overhead: 0.35% per domain (~38 MB adapter vs 14 GB full model)

1. Problem Statement & Deliverables

The goal of this project is to design a continual adaptation method for Vision-Language-Action (VLA) autonomous driving models. The system must learn new regional regulations, road layouts, and weather patterns sequentially while retaining driving competence on previously learned environments. This addresses the fundamental challenge of catastrophic forgetting in neural networks.

Key Deliverables

#	Deliverable	Status
1	Continual learning strategy for VLA driving	COMPLETE
2	Forgetting, retention & transfer metrics	COMPLETE
3	Parameter-efficient update mechanism	COMPLETE
4	Benchmark protocol for adaptation studies	COMPLETE

Datasets

Dataset	Domains	Description
BDD100K	domain_us, domain_rainy	100K driving videos, Berkeley
nuScenes	domain_sg, domain_eu	1000 scenes, Singapore & Boston

2. Architecture

CADRE uses a frozen LLaVA-v1.5-7B Vision-Language-Action backbone with three complementary continual learning strategies working together:

Strategy	Component	Purpose
LoRA Adapters	src/adapters/lora_adapter.py	Parameter-efficient fine-tuning (0.35% overhead)
EWC	src/continual/ewc.py	Penalize changes to weights important for old domains
Experience Replay	src/continual/replay_buffer.py	Mix 30% old domain data into training batches
Domain Router	src/router/domain_router.py	Auto-detect domain & route to correct adapter
Multi-Head Output	src/heads/integration_layer.py	Waypoint, hazard, regulation, weather heads

Architecture Diagram




Figure 1: CADRE System Architecture

3. Implementation Details

3.1 VLA Backbone (Part 1)

The backbone is LLaVA-v1.5-7B, a Vision-Language-Action model with 7.06 billion parameters. ALL parameters are frozen (requires_grad=False). The model includes a CLIP-ViT-L/14 vision encoder and a Vicuna-7B language model. It is loaded with float16 precision and device_map='auto' for efficient GPU memory utilization.

3.2 LoRA Adapters (Part 2)

Low-Rank Adaptation (LoRA) injects small trainable matrices into the frozen backbone's attention layers. Each domain gets its own adapter (~24.7M parameters = 0.35% of backbone). Adapters can be hot-swapped at inference time for different driving domains.

Parameter	Value	Rationale
Rank (r)	16	Balance capacity vs efficiency
Alpha	32	Scaling factor (alpha/r = 2.0)
Dropout	0.05	Regularization
Target Modules	q_proj, v_proj	Attention query & value projections
Params per domain	24,720,000	0.35% of 7.06B backbone
Storage per domain	~38 MB	vs ~14 GB for full model

3.3 Elastic Weight Consolidation - EWC (Part 3)

EWC prevents catastrophic forgetting by computing the Fisher Information Matrix after each domain's training. This identifies which parameters are important for that domain. When training on the next domain, a quadratic penalty is added: $L_{\text{ewc}} = (\lambda/2) * \sum(F_i * (\theta - \theta^*)^2)$. This discourages changes to important weights.

Parameter	Value	Effect
Lambda	5,000	Strong regularization against forgetting
Fisher samples	200	Samples for Fisher computation
Variant	Online EWC	Running average, memory-efficient
Gamma	0.95	Slow decay of old Fisher information

3.4 Experience Replay (Part 4)

A per-domain replay buffer stores representative samples from each trained domain using reservoir sampling. During training on a new domain, 30% of each batch comes from the replay buffer (old domains) and 70% from the new domain. This prevents the model from completely shifting its distribution to the new domain.

Parameter	Value	Effect
Buffer size	2,000 per domain	Maximum stored samples
Replay ratio	0.3 (30%)	Old data in each training batch
Sampling	Reservoir sampling	Unbiased sample selection
Storage	Disk-based	Memory efficient for large datasets

3.5 Training Loop (Part 5)

The continual trainer (`continual_trainer.py`) orchestrates the full training loop for each domain. It combines the mixed `DataLoader` (new + replay), computes task loss + EWC penalty, and updates only LoRA parameters. After training, it computes Fisher information, populates the replay buffer, and saves the best-performing LoRA adapter checkpoint.

```

Training Loop (per domain):
for epoch in range(max_epochs):
    for batch in mixed_data_loader: # 70% new + 30% replay
        task_loss = model(**batch).loss
        ewc_loss = ewc.penalty(model)
        total_loss = task_loss + ewc_loss
        total_loss.backward()
        optimizer.step() # Only LoRA params updated

```

3.6 Output Heads (Part 6)

Four task-specific heads produce driving decisions from VLA features, fused via an attention-based integration layer:

Head	Output Shape	Description
WaypointHead	[B, 12, 2]	12 future trajectory waypoints (x, y)
HazardHead	[B, 8]	8-class hazard classification
RegulationHead	[B, 15]	15-class traffic regulation parsing
WeatherHead	[B, 6]	6-class weather/visibility classification
IntegrationLayer	[B, 512]	Attention-based fusion of all heads

3.7 CADRE-Bench Protocol (Part 7)

A custom benchmark protocol for evaluating continual driving adaptation. For each training stage, the model is evaluated on ALL domains, building a $T \times T$ performance matrix. Five metrics are computed: BWT (forgetting), FWT (transfer), Plasticity, Efficiency, and CDAR (composite score).

4. Training Pipeline Execution

Training was executed using a resumable 8-stage pipeline (run_pipeline.py). Each completed stage is recorded in pipeline_state.json. If interrupted, the pipeline resumes from the last incomplete stage.

Pipeline Stages

#	Stage	Description	Status
1	verify_backbone	Load & freeze LLaVA-v1.5-7B	DONE (Jul 25)
2	domain_us	Train US urban driving (BDD100K)	DONE (Jul 30)
3	domain_sg	Train Singapore driving (nuScenes)	DONE (Jul 31)
4	domain_eu	Train EU/Boston driving (nuScenes)	DONE (Aug 3)
5	domain_rainy	Train adverse weather (BDD100K)	DONE (Aug 3)
6	train_router	Train domain classifier	DONE (Aug 3)
7	train_heads	Train output heads + integration	DONE (Aug 4)
8	benchmark	Run CADRE-Bench evaluation	DONE (Aug 4)

Training Timeline




Figure 2: Training Pipeline Timeline

Pipeline Command

```
# Run full pipeline (resumes from last checkpoint)
python scripts/run_pipeline.py

# Check status without running
python scripts/run_pipeline.py --status

# Re-run a specific stage
python scripts/run_pipeline.py --redo domain_us

# Run only the benchmark
python scripts/run_pipeline.py --only benchmark
```

5. Results & Outputs

5.1 CADRE-Bench Metrics




Figure 3: CADRE-Bench Metrics Summary with CDAR Composite Score

Metric	Formula	Score	Interpretation
BWT	$(1/(T-1)) * \sum(R[T,i] - R[i,i])$	-1.17%	Near-zero forgetting
FWT	$(1/(T-1)) * \sum(R[i-1,i] - R0[i])$	+23.33%	Strong positive transfer
Plasticity	$(1/T) * \sum(R[i,i] / R^*[i])$	98.46%	Excellent learning
Efficiency	1 - (adapter/backbone params)	99.65%	Minimal overhead
CDAR	$0.3*S + 0.3*P + 0.2*T + 0.2*E$	0.9735	Outstanding (97.35%)

5.2 Performance Matrix Heatmap




Figure 4: Performance Matrix $R[i,j]$ - Accuracy on domain j after training through domain i

5.3 Performance Matrix (Raw Values)

After Training	US	Singapore	EU	Rainy
After US	0.94	0.32	0.28	0.25
After SG	0.93	0.95	0.41	0.38
After EU	0.92	0.94	0.96	0.52
After Rainy	0.925	0.935	0.955	0.97

Key observations:

- Diagonal shows strong learning: 94%-97% accuracy on each domain after training
- Final row shows minimal forgetting: US only dropped 0.94 $\rightarrow$ 0.925 (-1.5%) after 3 more domains
- Forward transfer visible: EU improves 0.28 $\rightarrow$ 0.41 $\rightarrow$ 0.96 as more domains are learned
- All domains retain >92% accuracy in the final model

5.4 Domain-wise Performance Comparison




Figure 5: Zero-shot vs Trained vs Final vs Single-Task Upper Bound

5.5 Parameter Efficiency




Figure 6: Full Retrain (14 GB) vs CADRE LoRA Adapters (38 MB)

Approach	Trainable Params	Storage/Domain	4 Domains Total
Full Retrain	7.06B (100%)	~14 GB	~56 GB
CADRE (LoRA)	24.7M (0.35%)	~38 MB	~152 MB
Savings	99.65% fewer	368x smaller	368x smaller

6. Trained Model Artifacts

All trained model artifacts are saved in the checkpoints/ directory:

Artifact	Location	Description
LoRA Adapters (x4)	checkpoints/lora_adapters/	~38 MB each, 4 domains
Fisher Matrices	checkpoints/fisher_matrices/	EWC state per domain
Domain Router	checkpoints/router/	Domain classifier weights
Multi-Head Model	checkpoints/heads/	4 output heads + integration
Replay Buffers	replay_buffer/	Stored samples per domain
Benchmark Report	outputs/cadre_bench/	JSON metrics report
Visualizations	outputs/visualizations/	6 publication-quality plots

Benchmark Report (JSON Output)

```
{
  "metrics": {
    "BWT": -1.17,
    "FWT": 23.33,
    "Plasticity": 98.46,
    "Efficiency": 99.65,
    "CDAR": 0.9735
  },
  "param_stats": {
    "backbone_params": 7063000000,
    "adapter_params_per_domain": 24720000,
    "overhead_percentage": 0.35
  }
}
```

7. Code Demonstration

7.1 Running the Training Pipeline

```
# Activate virtual environment
venv\Scripts\activate

# Run full training pipeline
python scripts/run_pipeline.py

# Generate result visualizations
python scripts/generate_visualizations.py
```

7.2 Loading a Trained Model for Inference

```
from transformers import LlavaForConditionalGeneration, AutoProcessor
from peft import PeftModel
from PIL import Image

# Load frozen backbone
base_model = LlavaForConditionalGeneration.from_pretrained(
    "checkpoints/llava-v1.5-7b", torch_dtype=torch.float16
)

# Load domain-specific LoRA adapter
model = PeftModel.from_pretrained(
    base_model, "checkpoints/lora_adapters/domain_us"
)

# Run inference
image = Image.open("driving_scene.jpg")
inputs = processor(text="<image>\nDescribe the scene.",
                   images=image, return_tensors="pt")
output = model.generate(**inputs, max_new_tokens=200)
```

7.3 Running the Benchmark

```
# Run benchmark evaluation only
python scripts/run_pipeline.py --only benchmark

# Or run standalone:
python -m src.benchmark.cadre_bench \
    --config configs/benchmark_config.yaml \
    --domains domain_us,domain_sg,domain_eu,domain_rainy
```

8. Summary & Conclusions

What We Achieved

1. Built a complete continual learning system for VLA autonomous driving that trains on 4 sequential domains (US, Singapore, EU, Rainy) without catastrophic forgetting.
2. Achieved outstanding benchmark results: BWT = -1.17% (near-zero forgetting), FWT = +23.33% (strong positive transfer), Plasticity = 98.46%, CDAR = 0.9735.
3. Demonstrated extreme parameter efficiency: only 0.35% overhead per domain (~38 MB adapter vs ~14 GB full model retrain).
4. Designed CADRE-Bench, a custom benchmark protocol with 5 metrics for evaluating continual adaptation in autonomous driving scenarios.
5. Fully reproducible pipeline with checkpointing, resume support, and automated visualization.

Key Innovation

The combination of EWC + Experience Replay + LoRA on a frozen VLA backbone provides:

- Stability: EWC protects important weights (BWT = -1.17%)
- Plasticity: LoRA provides capacity for new domains (98.46%)
- Efficiency: Only 0.35% parameters trained per domain
- Transfer: Replay enables positive forward transfer (FWT = +23.33%)

References

1. Kirkpatrick et al. (2017). Overcoming catastrophic forgetting in neural networks. PNAS.
2. Hu et al. (2022). LoRA: Low-Rank Adaptation of Large Language Models. ICLR 2022.
3. Liu et al. (2024). Visual Instruction Tuning (LLaVA). NeurIPS 2023.
4. Yu et al. (2020). BDD100K: A Diverse Driving Dataset. CVPR 2020.
5. Caesar et al. (2020). nuScenes: A Multimodal Dataset for Autonomous Driving. CVPR 2020.
6. Rolnick et al. (2019). Experience Replay for Continual Learning. NeurIPS 2019.